

# CS5010: Lecture 9

## Intro to Concurrency

---

Fall 2022

Brian Cross, Northeastern University

Additional credits: Ian Gorton & Abi Evans

# Before we get started...

Clone the Code\_from\_Lectures repo from the course Github organization

[https://github.ccs.neu.edu/cs5010seaF22/Code\\_From\\_Lectures](https://github.ccs.neu.edu/cs5010seaF22/Code_From_Lectures)

Then, open the Lecture9 project in the Evening\_Lectures if you'd like to follow along.

# Administrivia

- Homework #4
  - Group project
  - Due Monday, November 7<sup>th</sup> by 11:59pm
  - Work in your team Group\_<id1>\_<id2> repo
- UML Draft Designs
- Codewalks 4 and 5
  - Combination of in-person and recorded
  - You'll want your team present for lectures to participate

# Agenda

---

- Concurrency overview
- Simple threads in Java
- Race conditions
- Deadlock
- Producer-consumer problem
- Thread states
- Thread-safe collections

# Concurrency



# Concurrency

## Doing multiple things at the same time

- Multiple applications
- Multiple processes within an application
- Distributed systems

Processes Performance App history Startup Users Details Services							
Name	Status	12% CPU	81% Memory	13% Disk	0% Network	0% GPU	GPU engine
Vmmem		3.9%	0 MB	0 MB/s	0 Mbps	0%	
Microsoft OneDrive (32 bit)		3.6%	313.9 MB	1.5 MB/s	0 Mbps	0%	
Microsoft Teams (4)		0.8%	74.5 MB	0.1 MB/s	0 Mbps	0%	
bzserv (32 bit)		0.6%	2.0 MB	0 MB/s	0 Mbps	0%	
Antimalware Service Executable		0.6%	100.9 MB	0 MB/s	0 Mbps	0%	
Task Manager		0.6%	24.1 MB	0 MB/s	0 Mbps	0%	
Service Host: Remote Desktop S...		0.4%	208.9 MB	0 MB/s	0 Mbps	0%	
System		0.3%	0.1 MB	0.1 MB/s	0 Mbps	0%	
IntelliJ IDEA (4)		0.2%	1,446.6 MB	0.1 MB/s	0 Mbps	0%	
Desktop Window Manager		0.2%	65.0 MB	0 MB/s	0 Mbps	0.3%	GPU 0 - 3D
System interrupts		0.1%	0 MB	0 MB/s	0 Mbps	0%	
Client Server Runtime Process		0.1%	0.7 MB	0 MB/s	0 Mbps	0%	
Windows Driver Foundation - U...		0.1%	159.3 MB	0 MB/s	0 Mbps	0.2%	GPU 0 - 3D
bzfilelist (32 bit)		0.1%	13.7 MB	0.2 MB/s	0 Mbps	0%	
PowerToys Runner		0.1%	1.0 MB	0 MB/s	0 Mbps	0%	
Service Host: Unistack Service G...		0.1%	4.3 MB	0.1 MB/s	0 Mbps	0%	
Service Host: Diagnostic Policy ...		0.1%	19.6 MB	0 MB/s	0 Mbps	0%	
RDP Clipboard Monitor		0%	1.4 MB	0 MB/s	0 Mbps	0%	
Windows Explorer (2)		0%	41.2 MB	0 MB/s	0 Mbps	0%	
PowerChute Data Service (32 bit)		0%	9.6 MB	0 MB/s	0 Mbps	0%	

# Concurrency

**Key: One task should not block or slow another**

- Tasks trigger/are triggered by events
- Event order is unpredictable
- Some tasks are time consuming e.g. reading large files
- System needs to stay *responsive*


Click Me!



# Processes and threads

Run a program → loads it in memory

- Contains one or more **process**
  - Independent: can't directly access data / resources in other processes






# Processes and threads

Run a program → loads it in memory

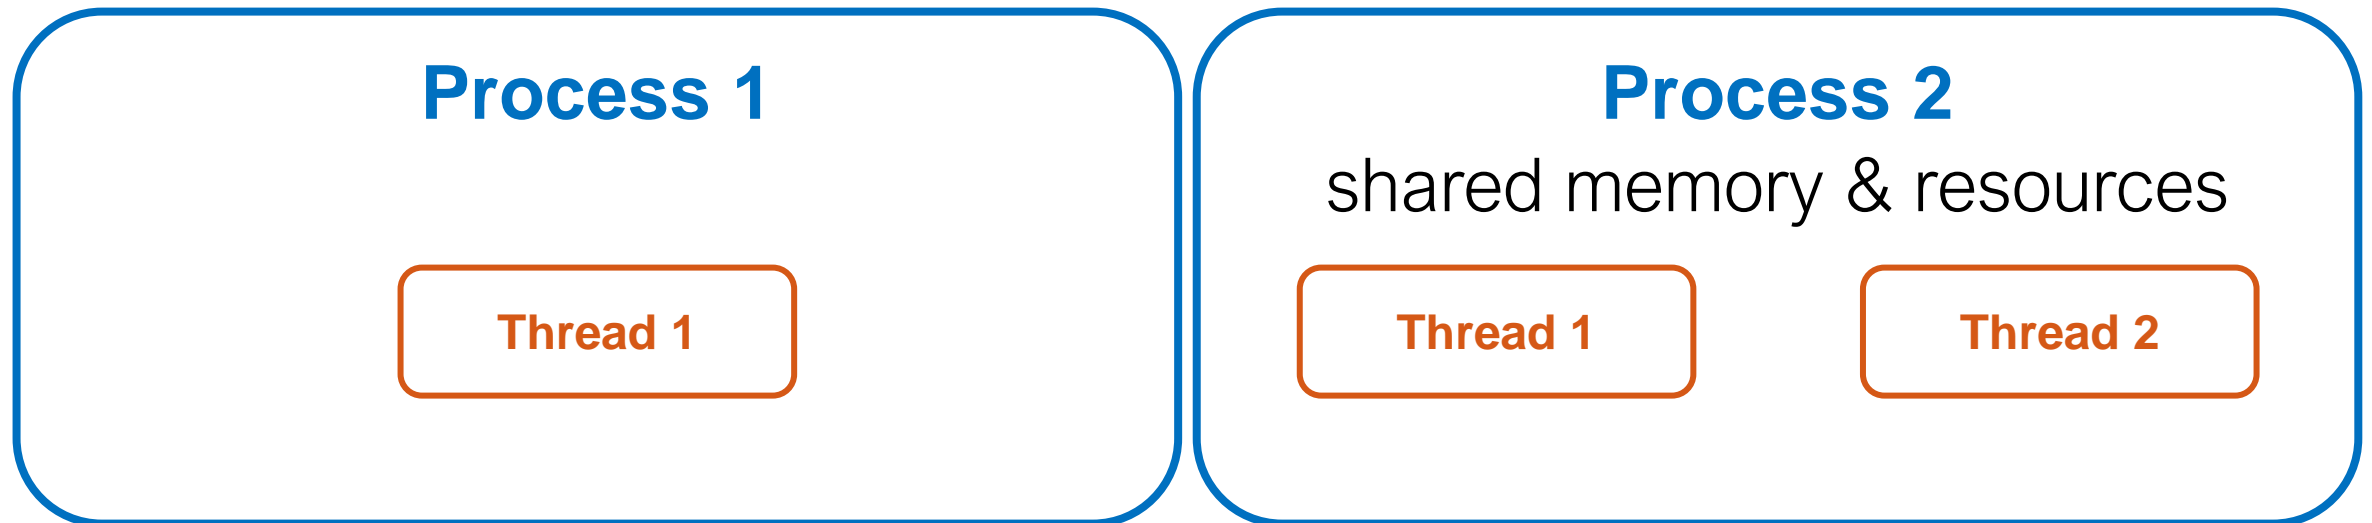
- Contains one or more **process**
- Each process has one or more **threads**



# Processes and threads











Run a program → loads it in memory

- Contains one or more **process**
- Each process has one or more **threads**
- In Java, you'll mostly work with threads



# Processes and threads example

All the processes and threads belonging to Firefox (with 8 tabs open)

Process Name	▼	% CPU	CPU Time	Threads	Idle Wake Ups	% GPU	GPU Time	PID	
followupd		0.0	0.51	2	0	0.0	0.00	757	at
fmfd		0.0	11.55	3	0	0.0	0.00	587	at
 FirefoxCP WebExtensions		0.0	2:29.65	36	1	0.0	0.00	24242	at
 FirefoxCP Web Content		0.1	3:10.92	43	10	0.0	0.00	56660	at
 FirefoxCP Web Content		1.5	4:19.17	44	15	0.0	0.00	57258	at
 FirefoxCP Web Content		0.1	1:02.62	41	2	0.0	0.00	56982	at
 FirefoxCP Web Content		0.0	29.48	39	2	0.0	0.00	57331	at
 FirefoxCP Web Content		0.0	2:49.35	46	2	0.0	0.00	57073	at
 FirefoxCP Web Content		0.6	1:46.70	43	3	0.0	0.00	57055	at
 FirefoxCP Web Content		0.1	29:42.15	50	3	0.0	0.00	24234	at
 FirefoxCP Web Content		0.2	6:59.58	50	78	0.0	0.00	25548	at
 Firefox		0.9	1:04:28.47	73	49	0.6	1:07:51.17	24232	at

Note the number of threads in each process!

# Threads

- **Lightweight** compared to processes
  - Threads within the same process share the same address space (memory)
  - Each thread has its own stack to support independent execution
- Communication between threads is easier than communication between processes
- Threads enable multi-tasking but shared resources mean its easy to make a mess!

# The main thread

Every application has at least one thread—the *main thread*

- You don't need to create it
- The main thread can create other threads
  - (These have to be created)

# Simple threads in Java

# Creating a thread in Java

Two options:

1. Create a **Runnable** object, pass it to a `java.lang.Thread` object
2. Subclass `java.lang.Thread`, which implements **Runnable**

Both options:

- Must override **Runnable**'s **run()** method (*where* differs for each approach)
- Call **Thread.start()** → calls the Runnable object's **run()** method

# Creating a thread by implementing Runnable

**Create a new class implementing interface Runnable...**

```
public class SimpleThread implements Runnable {  
    @Override  
    void run() {  
        // Do stuff when the thread is started e.g.  
        System.out.println("I'm a thread.");  
    }  
}
```



# Creating a thread by implementing Runnable

**Create a new class implementing interface Runnable...**

```
public class SimpleThread implements Runnable {  
    @Override  
    void run() {  
        // Do stuff when the thread is started e.g.  
        System.out.println("I'm a thread.");  
    }  
}
```

**The “driver” of the thread**

# Creating a thread by implementing Runnable

**...pass an instance of the Runnable object to Thread**

```
public class SomeClass {  
    public class void main() {  
        Thread t1 = new Thread(new SimpleThread());  
    }  
}
```

# Creating a thread by implementing Runnable

...pass an instance of the Runnable object to Thread

```
public class SomeClass {  
    public class void main() {  
        Thread t1 = new Thread(new SimpleThread());  
    }  
}
```

**Built-in class**

**Instance of an object  
which implements Runnable**

# Creating a thread by implementing Runnable

**Call `Thread.start()` to run the thread**

```
public class SomeClass {  
    public class void main() {  
        Thread t1 = new Thread(new SimpleThread());  
        t1.start();  
    }  
}
```

**Calls the `run()` method of the object passed to the `Thread` constructor**

In this example, `SimpleThread.run()`

# Creating a thread by subclassing Thread

**Create a class that extends Thread and override run()**

```
public class MyThread extends Thread {  
    @Override  
    void run() {  
        // Do stuff when the thread is started e.g.  
        System.out.println("I'm also a thread.");  
    }  
}
```

# Creating a thread by subclassing Thread

**Instantiate it and call `Thread.start()` to run the thread**

```
public class SomeClass {  
    public class void main() {  
        Thread t1 = new Thread(new SimpleThread());  
        t1.start();
```

```
        MyThread t2 = new MyThread();  
        t2.start();  
    }  
}
```

**Calls the `run()` method of the Thread subclass**

In this example, `MyThread.run()`

# Which option to choose?

## Implementing Runnable

- Slightly more code to write
- More flexible (Runnable object can still inherit another class)
- The best option most of the time

## Extending Thread

- Less code
- Not flexible (can't inherit any other class)
- Only for very simple applications

# Reasons why threads are tricky #1

When there are multiple threads, exact order of execution is not guaranteed or consistent → **interleaving**

```
public class RunnableThreadController {  
    public static void main(String[] args) {  
        System.out.println("This is the main thread starting. " + Thread.currentThread());  
        Thread t0 = new Thread(new BasicRunnable());  
        Thread t1 = new Thread(new BasicRunnable());  
        Thread t2 = new Thread(new BasicRunnable());  
        t0.start();  
        t1.start();  
        t2.start();  
        int x = 10;  
        int y = -5;  
        int z = x + y;  
        System.out.println("Time wasting... " + z);  
        System.out.println("This is the main thread finishing. " + Thread.currentThread());  
    }  
}
```



# Thread.sleep()

- Static method of the `java.lang.Thread` class
- Delays execution of the following code (within the thread)
- Must handle `InterruptedException`

# Thread.sleep() example

From sleepy.SleepyRunnable.java

```
@Override
public void run() {
    while (true) {
        System.out.println("Hello from " + Thread.currentThread().getName() + ". Time for a nap...");
        try {
            Thread.sleep( millis: this.nap_time * 1000);
        } catch (InterruptedException e) {
            System.out.println("I will stop napping then");
            return;
        }
        System.out.println(Thread.currentThread() + " is awake!");
    }
}
```

# Thread.sleep() example

From sleepy.SleepyRunnable.java

```
@Override
public void run() {
    while (true) {
        System.out.println("Hello from " + Thread.currentThread().getName() + ". Time for a nap...");
        try {
            Thread.sleep(millis: this.nap_time * 1000);
        } catch (InterruptedException e) {
            System.out.println("I will stop napping then");
            return;
        }
        System.out.println(Thread.currentThread() + " is awake!");
    }
}
```

Wait here for specified time...

...then execute what comes next

# interrupt()

Force a thread to stop what it's doing → throws an InterruptedException

```
Thread t = new Thread(new MyThread());  
t.start();  
t.interrupt();
```

# interrupt()

From sleepy.SleepyRunnable.java

```
@Override
public void run() {
    while (true) {
        System.out.println("Hello from " + Thread.currentThread().getName() + ". Time for a nap...");
        try {
            Thread.sleep( millis: this.nap_time * 1000);
        } catch (InterruptedException e) {
            System.out.println("I will stop napping then");
            return;
        }
        System.out.println(Thread.currentThread() + " is awake!");
    }
}
```

Execute if `interrupt()` called on an instance of `SleepyRunnable`



# Let's take a look

---

SleepyRunnable



# Joining threads with `thread.join()`

...forces one thread to wait until another thread completes its execution.

```
thread_to_wait_for.join();
```

```
Thread t1 = new Thread(new SomeThread());
```

```
Thread t2 = new Thread(new OtherThread());
```

```
Thread t3 = new Thread(new LastThread());
```

```
t1.start();
```

```
t2.start();
```


```
t2.join(); // will need to try-catch InterruptedException
```

```
t3.start(); // won't execute until t2 finishes
```

# Example use case for threads and sleep

Basic auto-saving text editor

- Automatically save the contents of the editor every X seconds.





# Basic auto-saving text editor

If we try using the main thread....

```
public class Controller {
    @FXML
    TextArea textArea;

    @FXML
    public void initialize() { this.saveText(); }

    private void saveText() {
        while (true) {
            try {
                Thread.sleep( millis: 5000);
                System.out.println(this.textArea.getParagraphs());
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}
```

**Everything** runs on the main thread

- (if we don't make new threads)

Sleeping this thread blocks anything from happening

# Basic auto-saving text editor

Moving the autosave operation to its own thread keeps the main thread responsive

```
public class Controller {  
    @FXML  
    TextArea textArea;  
  
    @FXML  
    public void initialize() {  
        this.startSaver();  
    }  
  
    private void startSaver() {  
        Thread saver = new Thread(new Saver(this.textArea));  
        saver.start();  
    }  
}
```

main thread

```
public class Saver implements Runnable {  
    private TextArea textArea;  
  
    public Saver(TextArea textArea) {  
        this.textArea = textArea;  
    }  
  
    @Override  
    public void run() {  
        System.out.println("starting");  
        while (true) {  
            try {  
                Thread.sleep(5000);  
                System.out.println(this.textArea.getParagraphs());  
            } catch (InterruptedException e) {  
                return;  
            }  
        }  
    }  
}
```

autosave thread

# Race conditions

# Reasons why threads are tricky #2

Interleaving can cause **race conditions**

→ When multiple threads share the same variables & change it at the same time

1. Read value from memory to register
2. Change value in register
3. Write register value back to memory:

thread 1:  $x = x + 6$ ;

thread 2:  $x = x + 1$

The result after N operations?

# Race conditions

Thread 1	Thread 2
Reads x into register	
Register value + 6	
Writes register value to x	
	Reads x into register
	Register value + 1
	Writes register value to x

Fine

Thread 1	Thread 2
Reads x into register	
Register value + 6	
	Reads x into register
	Register value + 1
	Writes register value to x
Writes register value to x	

Not fine

# Race conditions

- Same program, different results
  - Depends on how CPU schedules execution
  - Different interleavings produce different outcomes
- Extremely hard to debug
  - Not reproducible
  - Extremely unpleasant when they occur in production systems!

# Root cause: non-determinism

- Sequential programs exhibit deterministic behavior
  - We know what they going to do in which order
- Race conditions caused by **non-deterministic** behavior
  - We don't know exactly what's going to happen when
  - Non-determinism is not always a problem

# Two kinds of non-determinism

## Observable

Program may give different result  
(bad)

thread 1:  $x = x + 6$ ;  
thread 2:  $x = x + 1$ ;

## Non-observable

Program may execute differently,  
but result always the same (fine)

thread 1:  $a = 2$ ;  $b = a + 6$ ;  
thread 2:  $x = 9$ ;  $y = x - 3$

No shared variables → results are  
always the same



# Another race condition example

A ticketing system

- Multiple simultaneous ticket requests for the same seat  
→ if not properly handled, ticket could be oversold

**Take a look at Code\_from\_Lectures > Lecture9 > racecondition**

- **Can you see the problem?**

# Ticketing system example

In TicketRace:

```
public void takeSeat(String name) {  
    if (!this.isTaken) {  
        String message = "Taken by " + name;  
        System.out.println(message);  
        this.isTaken = true;  
        this.customerName = name;  
    }  
}
```

In Customer (thread):

```
@Override  
public void run() {  
    ticket.takeSeat(this.name);  
}
```

Customer 0: not taken

Customer 1: not taken

Customer 0: takes seat

Customer 1: takes seat

Seat taken by: customer 1  
(customer 0 is angry)

# Let's take a look

---

racecondition



# Common causes for race conditions

## Check-then-act

```
public void takeSeat(String name) {  
    if (!this.isTaken) {  
        String message = "Taken by " + name;  
        System.out.println(message);  
        this.isTaken = true;  
        this.customerName = name;  
    }  
}
```

# Common causes for race conditions

## Check-then-act

```
public void takeSeat(String name) {  
    if (!this.isTaken) {  
        String message = "Taken by " + name;  
        System.out.println(message);  
        this.isTaken = true;  
        this.customerName = name;  
    }  
}
```

## Read-modify-write

See Code\_from\_Lectures > Lecture 9 > **racecondition2**

- **Can you see the problem?**

# Read-modify-write example

```
public void incrementNumber() {  
    int num = this.number;  
    num++;  
    System.out.println("waste time");  
    this.number = num;  
}
```

- Thread X: reads number
- Thread X: wastes some time
- Thread Y: reads number
- Thread X: writes new number
- Thread Y: wastes some time
- Thread Y: writes new number starting from *old* value



# Let's take a look

---

racecondition2



# Why are race conditions hard to spot and debug?

- Both examples include extra code to slow processing down
- Simpler operations (e.g. `this.num++`) would work fine almost all of the time
- ...except when they don't

```
public void takeSeat(String name) {  
    if (!this.isTaken) {  
        String message = "Taken by " + name;  
        System.out.println(message);  
        this.isTaken = true;  
        this.customerName = name;  
    }  
}
```

```
public void incrementNumber() {  
    int num = this.number;  
    num++;  
    System.out.println("waste time");  
    this.number = num;  
}
```



# Avoiding race conditions

Synchronization, atomic data types

# Avoiding race conditions with locks

Use "locks" to impose ordering constraints

- Lock shared variables → can only be accessed by one thread at a time
- Each thread wishing to access a variable;
  - Takes the lock—everyone else has to wait
  - Changes the variable
  - Releases the lock—other threads can access

# Synchronized methods

**Add `synchronized` to the shared method definition:**

```
public synchronized void takeSeat(String name) {...}  
public synchronized void incrementNumber() {...}
```

- Prevents multiple invocations of synchronized methods on the same object from interleaving
- Handles taking and releasing the “lock”
- **Critical section:** a block of code that can't be accessed by more than one thread at a time

# Monitor locks

- Each Java object has a **monitor**, which a thread can lock or unlock
- Synchronized methods automatically lock the monitor on invocation by a thread → the thread has the lock
- Any other threads attempting to lock the monitor are blocked until the lock is available



# Let's take another look

---

Adding **synchronized**



# Synchronized: pros and cons

## Pro:

- Very simple to use

## Con:

- Can impact **liveness**—the ability of a program to do its work in a timely manner

Another option → Java's **atomic** data types

# Atomic Variables

- Defined in `java.util.concurrent.atomic`
- “Lock-free thread-safe” versions of single variable types

`Integer` → `AtomicInteger`

`Boolean` → `AtomicBoolean`

...and more

# Atomic Variables

Have unique methods...

```
private int number = 0;
```

```
→ private AtomicInteger number = new AtomicInteger(0);
```

```
this.number++;
```

```
→ this.number.incrementAndGet();
```

```
return this.number;
```

```
→ return this.number.get();
```



# Example: AtomicInteger

Run racecondition2.ReadModifyWrite.java

→ unpredictable results

Run atomicsolution.ReadModifyWrite.java


→ consistent results



# Let's take a look

---

atomicsolution



# Deadlock

# Reasons why threads are tricky #3

**Deadlock** occurs when multiple threads are waiting on each other → no thread can move forward.

Classic example: the dining philosophers problem


# The dining philosophers problem

5 philosophers, 5 bowls of noodles &  
5 chopsticks

Each philosopher must alternatively  
think and eat (forever).


- can take a left or right chopstick if available
- BUT can only eat if BOTH left and right chopsticks are available

**Goal:** ensure no-one starves!



# The dining philosophers problem

Philosopher behavior:



# Pseudo-code for a philosopher

```
while (true) {  
    think();  
    pick_up_left_chopstick();  
    pick_up_right_chopstick();  
    eat();  
    put_down_right_chopstick();  
    put_down_left_chopstick();  
}
```

# Deadlock

Run:

Code\_from\_Lectures > Lecture9 > deadlockedphilosophers\Table.java

**What's gone wrong?**



# Deadlock

2 threads sharing access to 2 shared variables via locks:


1. Thread 1: takes lock a
2. Thread 2: takes lock b
3. Thread 1: blocked from accessing b
4. Thread 2: blocked from accessing a




# Let's take a look

---

Deadlocked philosophers



# Deadlocked philosophers




# This is why concurrency is hard!

- Too few ordering constraints → race conditions
- Too many ordering constraints → deadlocks
- Hard/impossible to reason about based on modularity
  - Need to think about what all threads could do in all possible orders
- Thorough testing is impossible
  - Infinite number of possible interleavings
  - Controlled by system scheduler and events, not the program

# One solution


Number the chopsticks

Each philosopher picks up their lowest-numbered chopstick first




# One solution: numbered chopsticks

Gets there first





# One solution: numbered chopsticks




Lowest-numbered  
chopstick is taken  
→ waits

# One solution: numbered chopsticks

Everyone else can pick up their lowest-numbered chopstick...

Lowest-numbered chopstick is taken  
→ waits



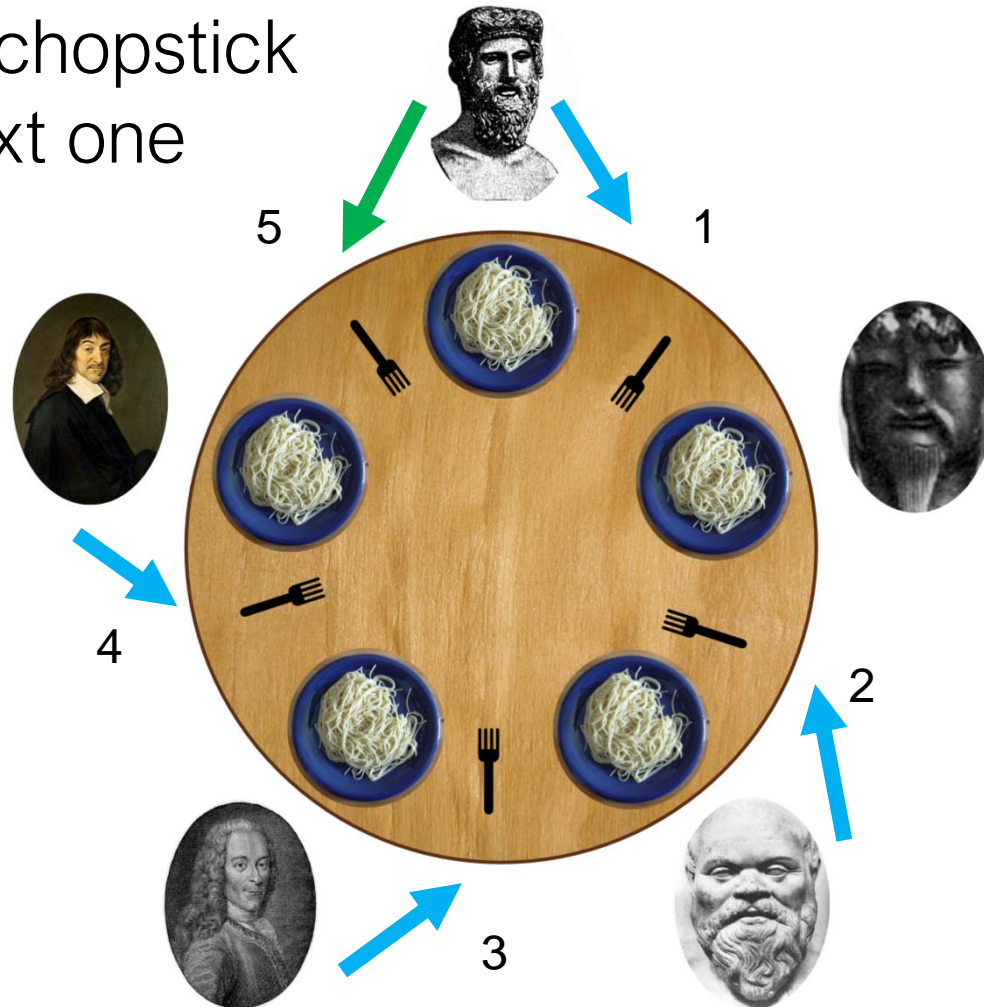


# One solution: numbered chopsticks

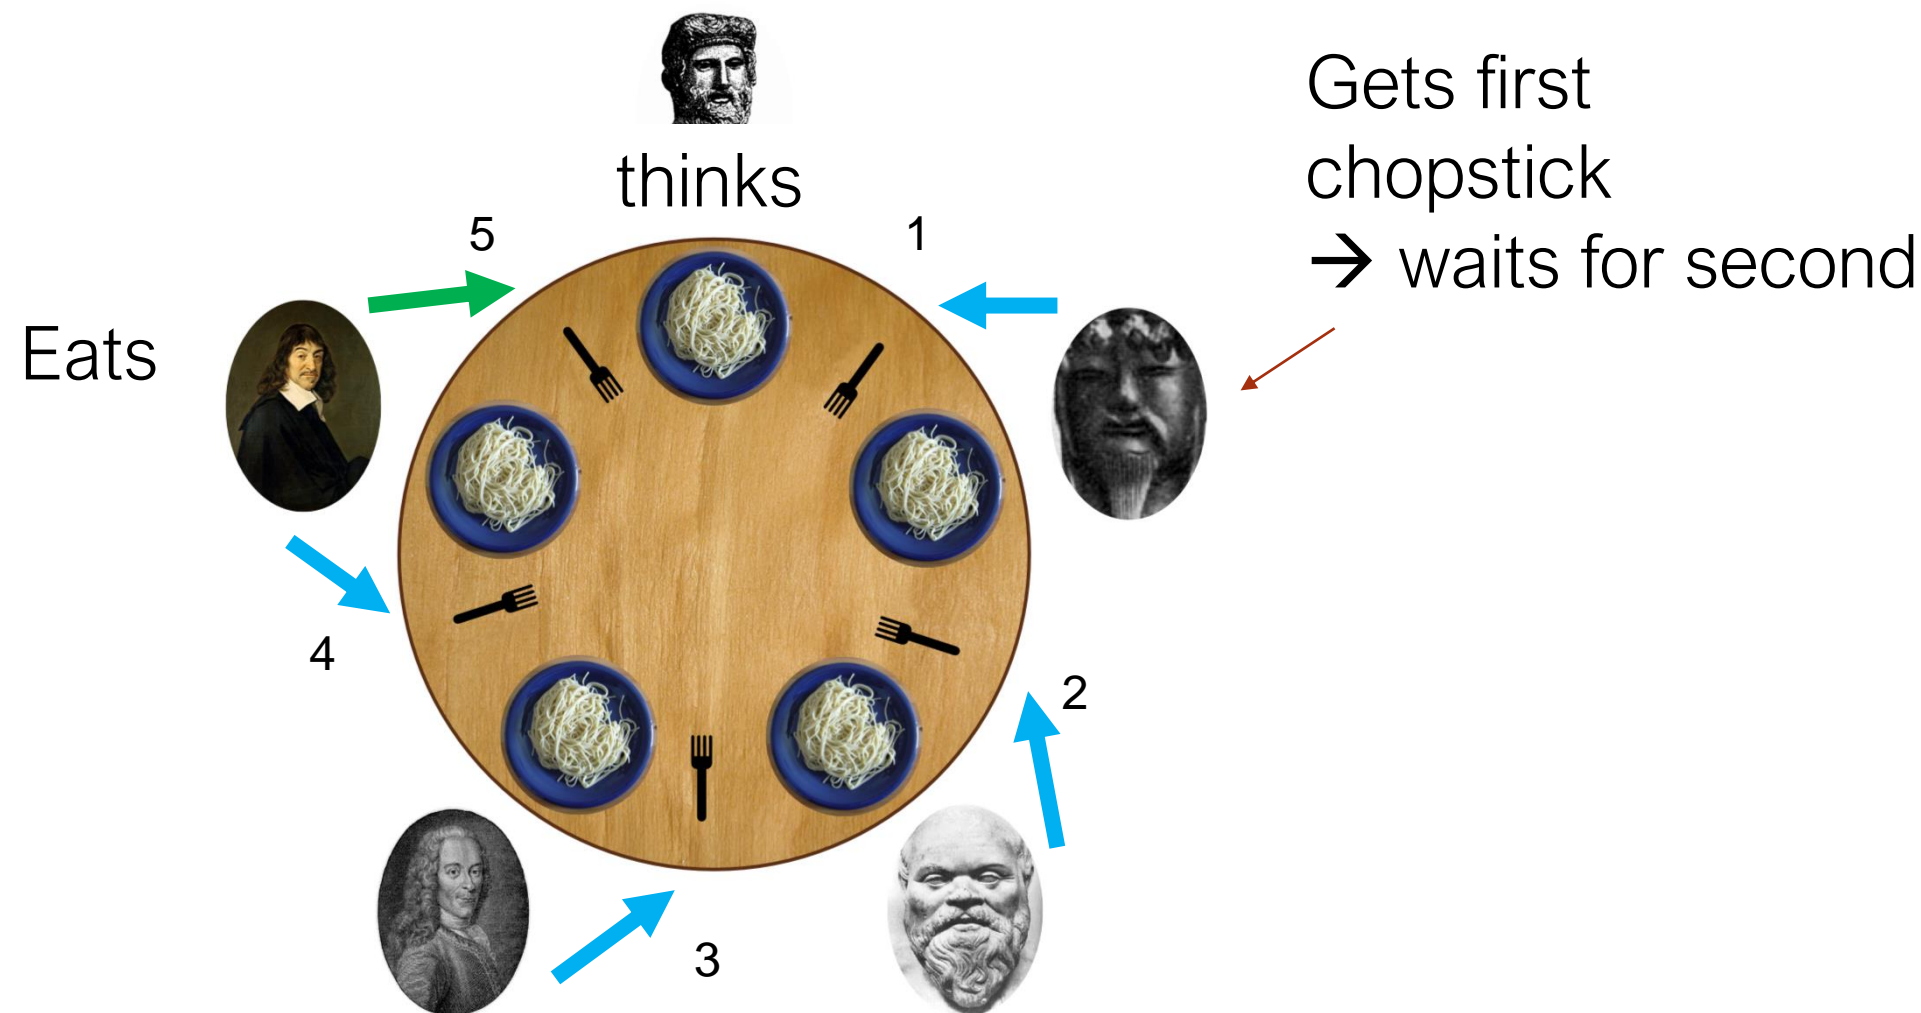
Philosophers with a chopstick try to pick up the next one

- All but one must wait

Lowest-numbered chopstick is taken  
→ waits



# One solution: numbered chopsticks





# Let's take another look


---

fed philosophers

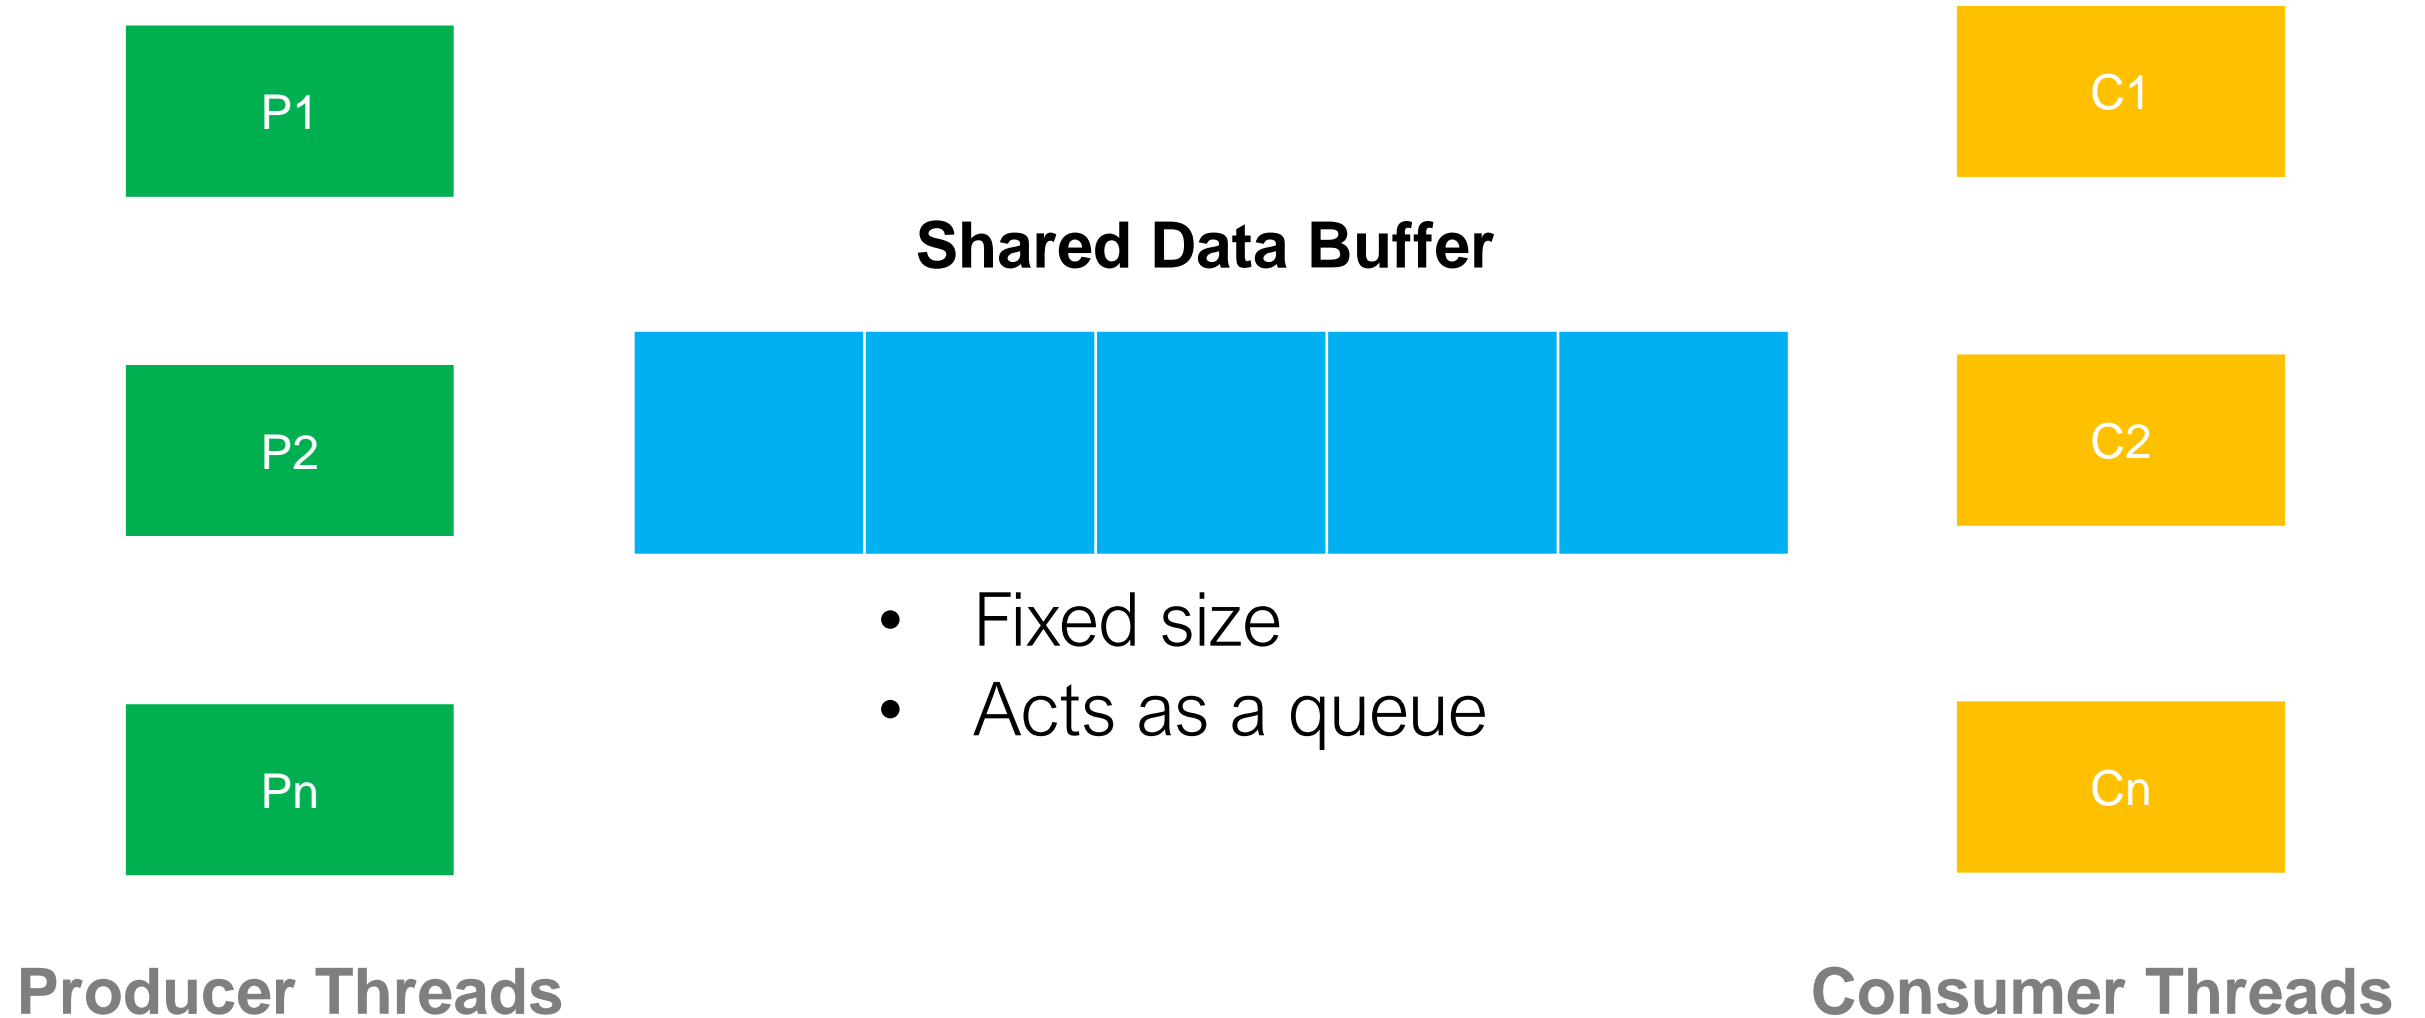
# Producer Consumer Problem



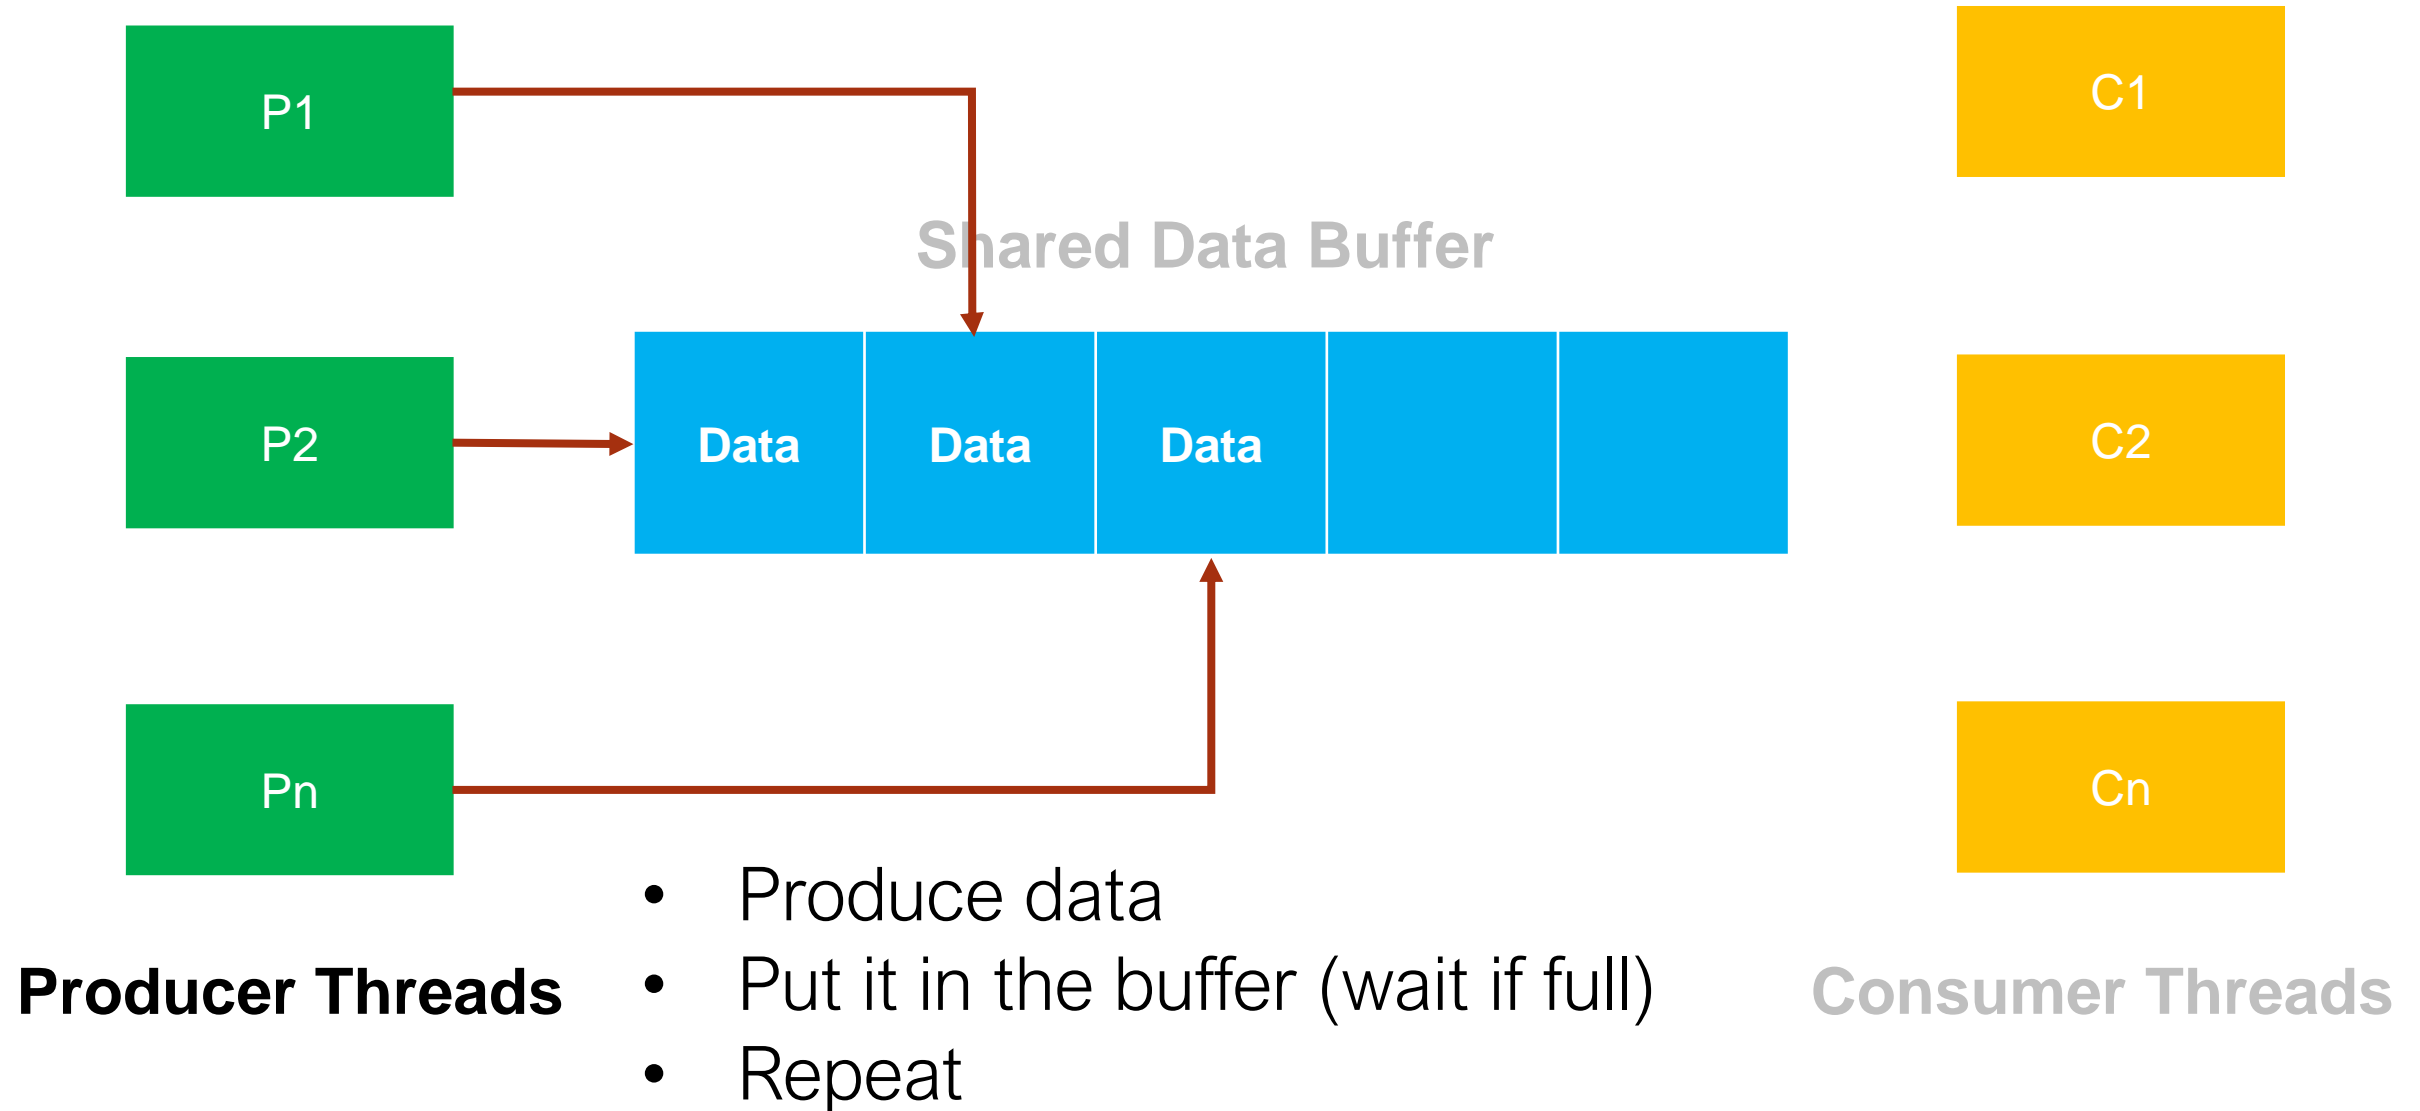
# Producers and consumers



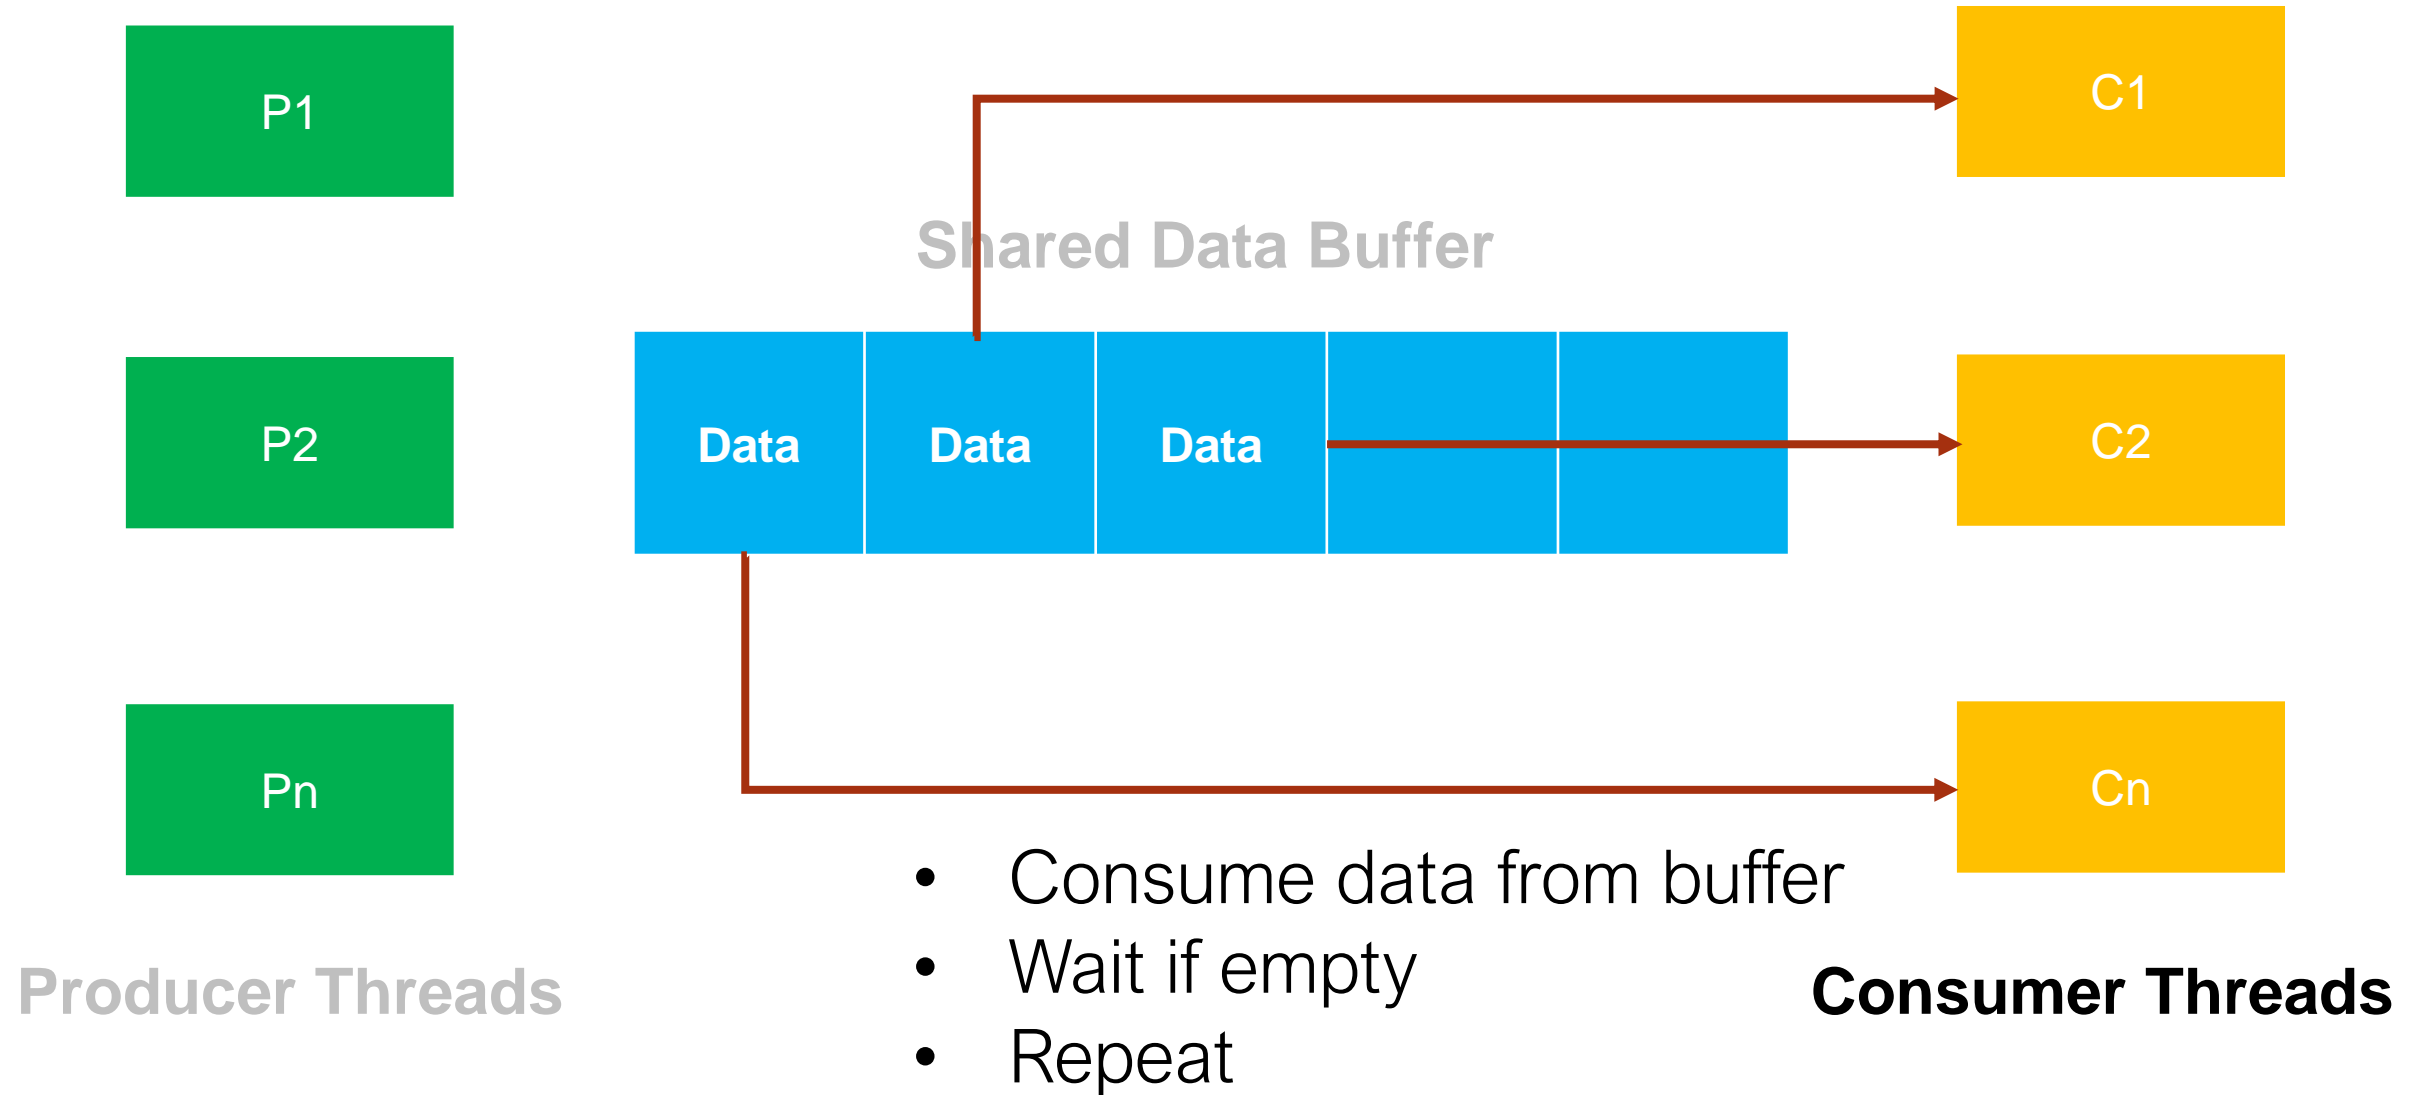
# Producers and consumers



# Producers and consumers

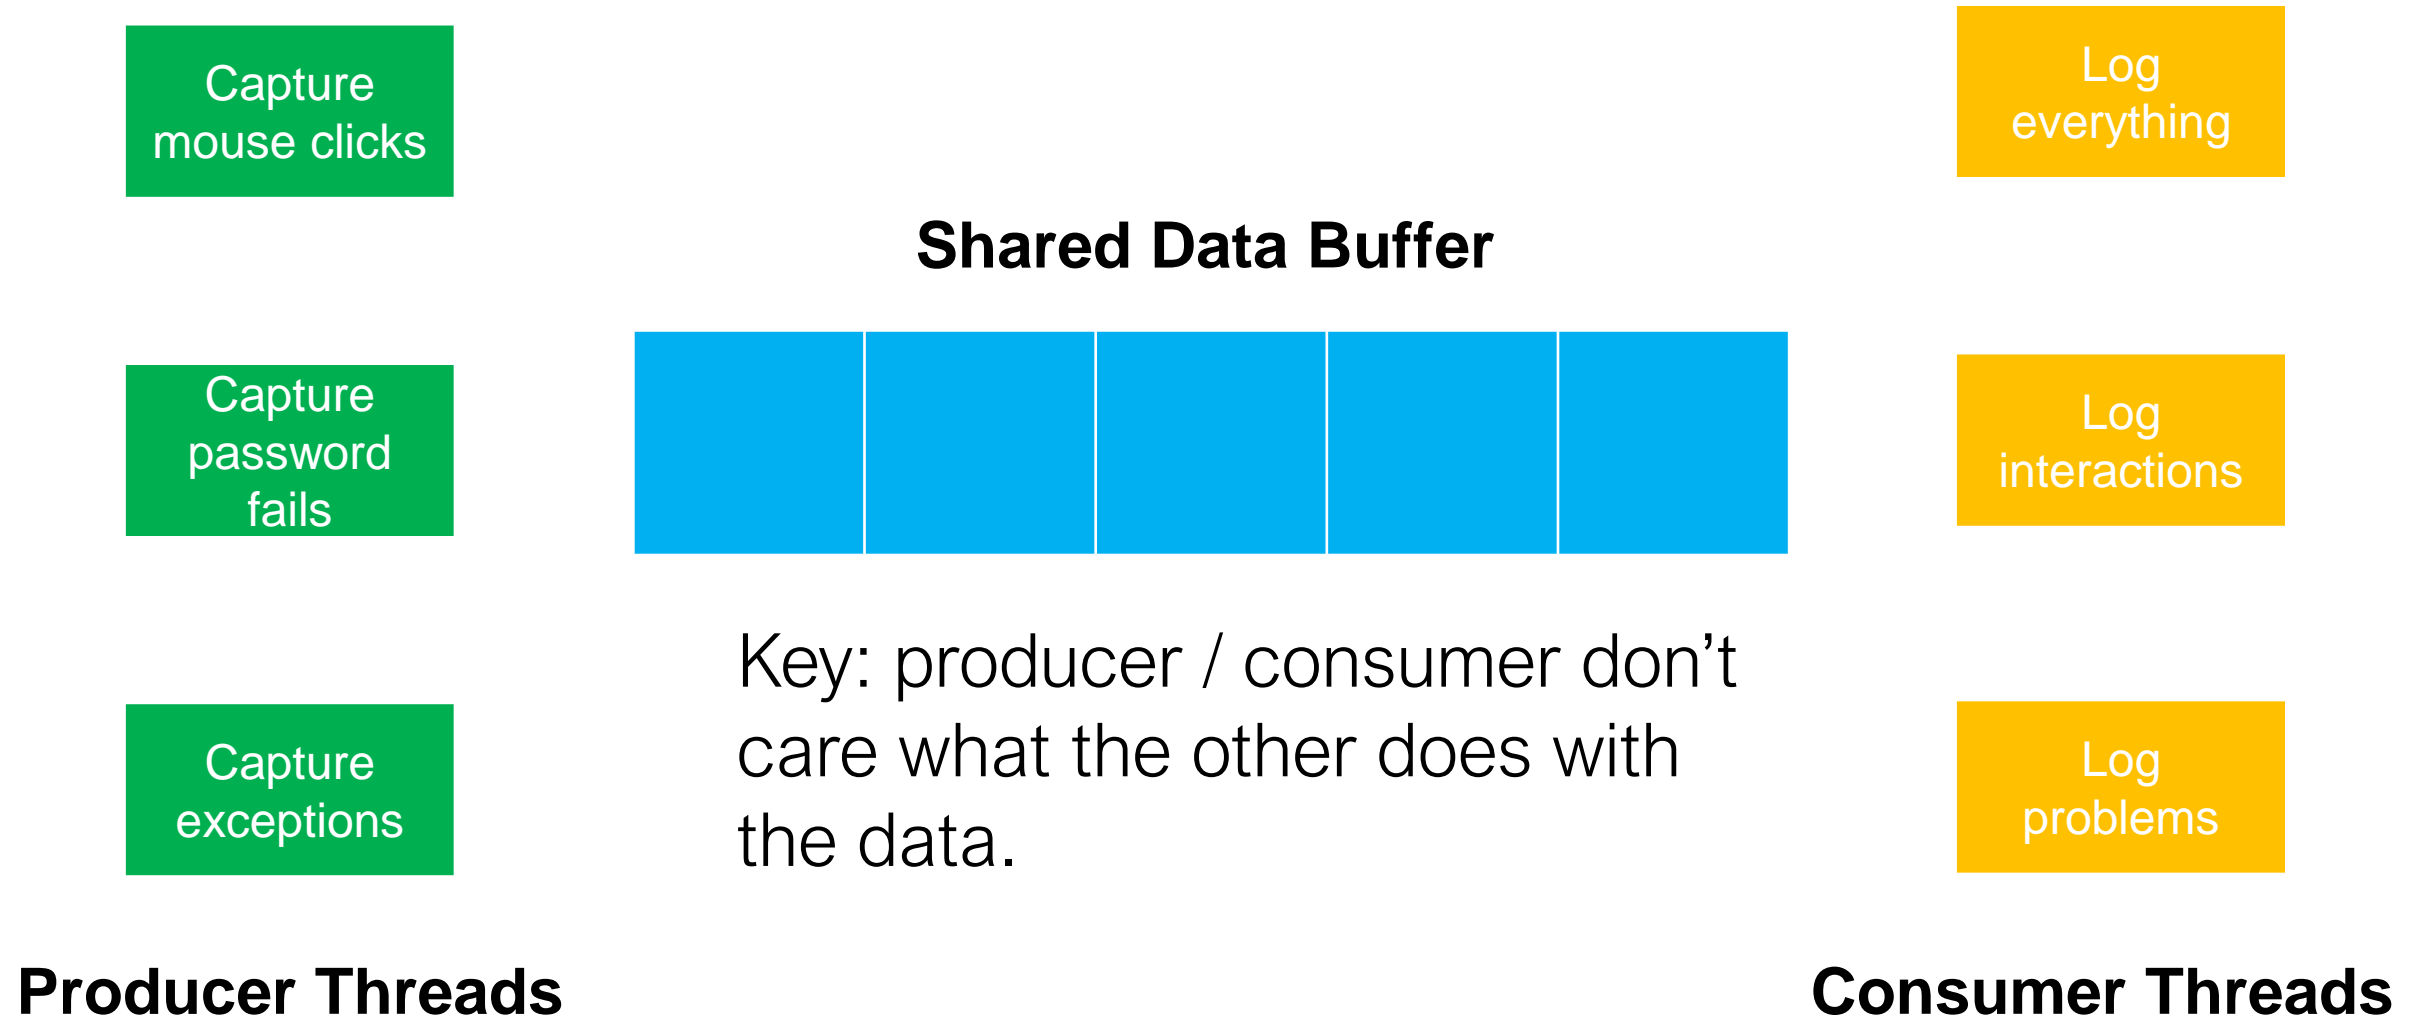


# Producers and consumers





# Use case: event logging



# Shared buffer implementation

- Some kind of data storage e.g. an integer, a list etc.
- Method to allow producer to add data
- Method to allow consumer to remove data
- Producer and consumer objects need reference to the shared buffer

# Producer Consumer: first try pseudocode

```
producer() {  
    while(true) {  
        item = produceItem();  
        if (bufferFull()) {  
            sleep();  
        }  
        putIntoBuffer(item)  
        if (!bufferEmpty()) {  
            wakeup(consumer);  
        }  
    }  
}
```

```
consumer() {  
    while(true) {  
        if (bufferEmpty()) {  
            sleep();  
        }  
        consumeItem()  
        if (!bufferFull()) {  
            wakeup(producer);  
        }  
    }  
}
```

# Potential for deadlock...

- Consumer notices empty buffer, moves inside the if block

# Potential for deadlock...

- Consumer notices empty buffer, moves inside the if block
- Consumer interrupted (context switch) before calling sleep

# Potential for deadlock...

- Consumer notices empty buffer, moves inside the if block
- Consumer interrupted (context switch) before calling sleep
- Producer creates an item, puts it in the buffer

# Potential for deadlock...

- Consumer notices empty buffer, moves inside the if block
- Consumer interrupted (context switch) before calling sleep
- Producer creates an item, puts it in the buffer
- Because the buffer was empty, producer tries to wake up consumer

# Potential for deadlock...

- Consumer notices empty buffer, moves inside the if block
- Consumer interrupted (context switch) before calling sleep
- But...consumer wasn't sleeping, wake up call is lost
- Producer creates an item, puts it in the buffer
- Because the buffer was empty, producer tries to wake up consumer



# Potential for deadlock...

- Consumer notices empty buffer, moves inside the if block
- Consumer interrupted (context switch) before calling sleep
- But...consumer wasn't sleeping, wake up call is lost
- Consumer goes to sleep
- Producer creates an item, puts it in the buffer
- Because the buffer was empty, producer tries to wake up consumer

# Potential for deadlock...

- Consumer notices empty buffer, moves inside the if block
- Consumer interrupted (context switch) before calling sleep
- But...consumer wasn't sleeping, wake up call is lost
- Consumer goes to sleep
- Producer creates an item, puts it in the buffer
- Because the buffer was empty, producer tries to wake up consumer
- Producer loops until the buffer fills, then goes to sleep

# Guards

Producer-Consumer style implementations require “guards” on the buffer

- Buffer accepts incoming data unless full → waits until room
- Buffer allows data to leave unless empty → waits until there is data

# Java guards (AKA monitors)

wait() and notify() statements → communicate to threads synchronized on the same object

E.g.:

- Thread A calls synchronized method of buffer
- Buffer isn't ready → wait() → thread A waits
- Thread B changes buffer using different method
- Buffer is ready to complete thread A's request → notifyAll() → all threads waiting on method continue

Example: Code\_from\_Lectures > Lecture9 > producerconsumerguarded



# Let's take a look

---

Producer consumer guarded



# Takeaways: the shared buffer

(This example shows one approach of several)

Buffer implemented in its own class:

- Has a data structure for storing things sent from/to threads
- Data structure has a size limit
- Methods for putting data and getting data are `synchronized`
- General approach to putting and getting:
  - if buffer is full/empty  $\rightarrow$  `wait()`
  - do job
  - `notifyAll()`

# Takeaways: producer and consumer

Producer and consumer are both threads (implementing Runnable or subclassing Threads)

- Each has a reference to the same buffer instance
  - e.g. if a custom class: pass it via the constructor
- The `run()` method starts the process of putting/getting
  - All checking/waiting handled by the buffer...no need to do it here

# Thread states



# Thread states

- New thread
  - Created but not started
- Runnable state
  - Started and running or ready to run
- Blocked/waiting state
  - Temporarily inactive due to another thread
- Dead state
  - `run()` terminates

# Blocked/waiting state

A thread is not runnable if one of the following occurs:

- `sleep()` is invoked
- `suspend()` is invoked
- `wait()` is invoked
  - waits for notification of a free resource
  - waits for completion of another thread
  - waits to acquire a lock on an object
- blocked on an I/O request

# Thread resumption

- If a thread is asleep:
  - Sleep period elapses or it is interrupted
- If a thread is suspended:
  - Its resume() method must be called
- If a thread is waiting:
  - The object owning the shared resource must relinquish it by calling either notify() or notifyAll()
- If waiting on I/O, I/O must complete

# Thread priority

- Every Java thread has a priority
  - High priority threads get scheduled more frequently than lower priority threads
- Java threads inherit priority from parent...
  - MIN\_PRIORITY (1)
  - NORM\_PRIORITY (5) Default
  - MAX\_PRIORITY (10)
- ...unless set using `thread.setPriority(priority)`

# Thread scheduling

- Scheduler chooses the runnable thread with the highest priority
- When there are multiple threads to choose from → picks one. The chosen thread runs until:
  - a higher priority thread becomes runnable
  - it yields or completes
  - its time allotment has expired

# Reentrancy

Every Java object has a lock associated with it

- Known as the intrinsic lock
- AKA *monitor* or *mutex* locks

Synchronized methods exploit this intrinsic lock

- Lock acquired by executing thread before entering a synchronized block
- Lock released automatically when the thread exits the synchronized block

# Reentrancy

Intrinsic locks are **reentrant**:

- If a thread tries to acquire a lock it already holds, it succeeds

Reentrancy facilitates encapsulation of locking behavior and simplifies OO concurrent code

# Thread safety

- Thread safety requires the internal state of an object to be protected from concurrent updates
  - i.e. can't be changed by more than one thing at a time
- What if an object has no state that persists...?
- Or cannot be modified by a calling thread?
- Is this thread-safe?



# Thread safety

**Stateless and immutable objects are always thread-safe**

# Thread safe collections


# Thread-safe collection classes

- Standard collection class are NOT thread-safe
  - i.e. multiple threads can modify a standard collection at the same time
- `java.util.concurrent` package provides thread-safe collections, including:
  - **BlockingQueue**: queue that blocks or times out when you attempt to add to a full queue, or retrieve from an empty queue
  - **ConcurrentMap**: thread-safe subinterface of `java.util.Map`
  - **ConcurrentHashMap**: thread-safe version of `HashMap`

# ConcurrentHashMap


A HashMap is divided into buckets

Lock applied at the bucket level



# BlockingQueue

Saves the buffer from having to wait() and notifyAll()



Produce data  
Put it in the queue (wait if full)  
Repeat

Wait if empty  
Retrieve from queue  
Repeat

# Using a BlockingQueue

All producer/consumer threads that will interact with the shared queue need a reference to it. E.g.:

```
public class Producer implements Runnable {  
    private BlockingQueue<Integer> buffer;  
    public Producer(BlockingQueue<Integer> buffer) {  
        this.buffer = buffer;  
    }  
    ... produce things...  
}
```

# Using a BlockingQueue

All producer/consumer threads that will interact with the shared queue need a reference to it. E.g.:

```
public class Consumer implements Runnable {  
    private BlockingQueue<Integer> buffer;  
    public Consumer(BlockingQueue<Integer> buffer) {  
        this.buffer = buffer;  
    }  
    ... consume things...  
}
```

# Using a BlockingQueue

Producers and consumers of the same BlockingQueue need a reference to the *same* instance...e.g.:

```
public class Driver {  
    public static void main(String[] args) {  
        ArrayBlockingQueue<Integer> shared_buffer  
            = new ArrayBlockingQueue<Integer>(10);  
        Thread p = new Thread(new Producer(shared_buffer));  
        Thread c = new Thread(new Consumer(shared_buffer));  
    }  
}
```



# Takeaways: BlockingQueue example

(Again, one of many possible configurations)

- No need for a separate buffer class
  - No need for:
    - synchronized put/get methods
    - wait()
    - notifyAll()
- BlockingQueue handles everything



# Let's take a look

---

Producer consumer  
blockingqueue

# Recap

- Concurrency is fundamental to software systems
- Introduces problems of race conditions and deadlocks
- Synchronization required as a solution
- Threads move through various states during the lifetime
- Scheduler makes decisions on which thread to run based on state and priority



# Questions?

---



# Implementation exercise



# Producer-consumer: file I/O

Use a producer-consumer approach to read a csv file and print it to the terminal. Try at least one of the following:

- Using a guarded buffer (`synchronized`, `wait()`, and `notifyAll()`)
- Using a thread-safe collection e.g. `BlockingQueue`