

Lecture 12: Functional Programming


Spring 2022


Instructor: Brian Cross

Lots of credit: Divya Chaudhary, Adrienne Slaughter



Lecture Agenda

- Functional Programming
 - Programming Paradigms
 - Motivation
 - Terminology
- 



Introduction

Terminology

- procedural programming
- object-oriented programming
- generic programming
- functional programming
- declarative programming
- imperative programming
- stream
- lambda, lambda expression
- immutability
- concurrency
- reduction
- external vs internal iteration
- terminal operation
- arrow token
- lazy evaluation
- eager
- method reference
- infinite streams

Functional Programming: What we do?

1

Start with a **stream** of data (primitive or objects)

2

Apply a series of operations or transformations to the stream

3

Reduce the stream to a single number or **collect** the stream to collection



So many
questions...

What's a stream, and
is a list a stream?

An array?

A hashmap?

How many times have you written code like this?

```
List<Record> records = new ArrayList<>();  
  
int total = 0;  
  
for (int i=0; i<records.size(); i++){  
    total += records.get(i).value();  
}
```

How many times have you written code like this?

```
List<Record> records = new ArrayList<>();  
  
int total = 0;  
  
for (int i=0; i<records.size(); i++) {  
    total += records.get(i).value();  
}
```

What could go wrong?

How many times have you written code like this?

```
List<Record> records = new ArrayList<>();  
  
int total = 0;  
  
for (int i=0; i<records.size(); i++){  
    total += records.get(i).value();  
}
```

External Iteration:

The programmer specifies the iteration details.

Let's simplify for a moment.

```
int total = 0;  
  
for (int i=0; i<10; i++){  
    total += i;  
}
```

Let's simplify for a moment.

```
int total = 0;

for (int i=0; i<10; i++){
    total += i;
}
```

```
int total = IntStream.rangeClosed(1, 10)
                      .sum();
```

Let's simplify for a moment.

```
int total = 0;

for (int i=0; i<10; i++){
    total += i;
}
```

```
int total = IntStream.rangeClosed(1, 10)
                      .sum();
```

“For the stream of ints from 1 to 10, calculate the sum.”

Stream and Stream Pipeline

- Stream: sequence of elements
- Stream pipeline: sequence of tasks (“processing steps”) applied to elements of a stream
- A stream starts with a data source.
 - Examples:
 - Terminal I/O
 - Socket I/O
 - File I/O
- A stream can generally be used like a queue— you’re reading from it, but you can’t go back in the stream. Once you’ve pulled an element off the stream, it’s no longer in the stream.

The stream

```
int total = IntStream.rangeClosed(1, 10)
                    .sum();
```

`IntStream` produces a stream of integers in the given range. `rangeClosed` is closed— produces ints including 1 and 10.

The Stream Pipeline

```
int total = IntStream.rangeClosed(1, 10)  
                        .sum();
```

The processing step to take, or task to complete using the stream.

The Stream Pipeline

```
int total = IntStream.rangeClosed(1, 10)  
                      .sum();
```

The processing step to take, or task to complete using the stream.

Reduction:

Reduces the stream of values into a single value.

The Stream Pipeline

```
int total = IntStream.rangeClosed(1, 10)  
                      .sum();
```

The processing step to take, or task to complete using the stream.

Internal Iteration:

`IntStream` handles all the iteration details— we don't write them ourselves.

Reduction:

Reduces the stream of values into a single value.

The Stream Pipeline

Declarative Programming:

Internal Iteration:

`IntStream` handles all the iteration details— we don't write them ourselves.

Imperative Programming:

External Iteration:

The programmer specifies the iteration details.

The Stream Pipeline

Declarative

Programming:

Specifies *what* to do

Internal Iteration:

`IntStream` handles all the iteration details— we don't write them ourselves.

Imperative

Programming:

Specifies *how* to do something.

External Iteration:

The programmer specifies the iteration details.

The Stream Pipeline

```
int total = IntStream.rangeClosed(1, 10)
                      .sum();
```

intRange()



But what if we want to sum the even numbers between 2 and 20?

sum()

Summing even ints 2-20

```
int total = IntStream.rangeClosed(1, 10)
    .map((int x) -> {return x * 2;})
    .sum();
```

This converts the stream from 1:10 to 2:20 by multiplying by 2.

intRange()



map()



sum()

.map()

- Takes a method, and applies it to every element in the stream.

```
.map((int x) -> {return x * 2;})
```

Wait, what? A *method*?

lambdas: anonymous methods

- *lambda* or *lambda expression*
 - aka *anonymous method*
 - aka method-without-a-name
 - aka the method that shall not be named

```
(int x) -> {return x * 2;}
```

lambdas: anonymous methods

- Methods that can be treated as data
 - pass lambdas as arguments to other methods (map)
 - assign lambdas to variables for later use
 - return a lambda from a method

```
(int x) -> {return x * 2;}
```

lambdas: syntax

```
(parameter list) -> {statements}
```

```
(int x) -> {return x * 2;
```

Parameter: one `int` named `x`

Statement: `return 2*x`

lambdas: syntax

```
(parameter list) -> {statements}
```

```
(int x) -> {return x * 2;
```

Same as:

```
int multiplyBy2(int x) {  
    return x * 2;  
}
```


Difference:

- the lambda doesn't have a name
- compiler infers return type

lambdas: simplifying syntax

Eliminate parameter type

```
(int x) -> {return x * 2;}
```




```
(x) -> {return x * 2;}
```

Type is inferred.
If it can't be inferred,
compiler throws an
error.

lambdas: simplifying syntax

Simplify the body

```
(x) -> {return x * 2;}
```




```
(x) -> x * 2
```

- return is inferred
- semicolon and brackets not necessary

lambdas: simplifying syntax

Simplify parameter list

`(x) -> x * 2`



`x -> x * 2`

Can remove parens for
single parameter


lambdas: simplifying syntax

lambda with no parameters

```
() -> System.out.println("Hello Lambda!")
```

Method Referenes

```
.map(x -> x.toUpperCase())
```



```
.map(String::toUpperCase)
```

```
objectName::instanceMethodName
```


Sometimes, you want to just pass the incoming parameter to another method.

lambdas: scope

- Lambdas do not have their own scope
 - Can't shadow a method's local variable with lambda params with the same name
 - Lambdas share scope with the enclosing method



Fun Note:

- Why is it called lambda?
 - Alonzo Church: Church's λ calculus
 - Developed to provide a rigorous foundation for studying functions and function application
- 

Stream Pipeline: Intermediate & Terminal Operations

```
int total = IntStream.rangeClosed(1, 10)
                      .map((int x) -> {return x * 2;})
                      .sum();
```

- `map()` is an *intermediate* operation
- `sum()` is a *terminal* operation

Stream Pipeline: Intermediate & Terminal Operations

```
int total = IntStream.rangeClosed(1, 10)
                      .map((int x) -> {return x * 2;})
                      .sum();
```

- `map()` is an *intermediate* operation
- `sum()` is a *terminal* operation

Intermediate operations use *lazy evaluation*.

The operation produces a new stream object, but no operations are performed on the elements until the terminal operation is called to produce a result.

Stream Pipeline: Intermediate & Terminal Operations

```
int total = IntStream.rangeClosed(1, 10)
                      .map((int x) -> {return x * 2;})
                      .sum();
```

- `map()` is an *intermediate* operation
- `sum()` is a *terminal* operation

Terminal operations are *eager*.
The operation is performed when called.

Examples

Intermediate Operations

- `filter()`
- `distinct()`
- `limit()`
- `map()`
- `sorted()`

Terminal Operations

`forEach()`

`collect()`

Reductions:

- `average()`
- `count()`
- `max()`
- `min()`
- `reduce()`

Back to our example...

```
int total = IntStream.rangeClosed(1, 10)
                      .map((int x) -> {return x * 2;})
                      .sum();
```

For this example, we chose to create a stream of even ints from 2 to 20 by mapping from 1:10, multiplying by 2.

How else can we do this?

Back to our example...

```
int total = IntStream.rangeClosed(1, 20)
                    .filter(x -> x%2 == 0)
                    .sum();
```

Filter!

The lambda for the filter operation needs to return a boolean indicating whether the given element should be in the output stream.

Clarifying elements through the pipeline

```
int total = IntStream.rangeClosed(1, 10)
    .filter(
        x -> {
            System.out.printf("%nFilter: %d%n", x);
            return x % 2 == 0;
        })
    .map(
        x -> {
            System.out.printf("map: %d", x);
            return x * 3;
        })
    .sum();
System.out.println("\n\nTotal: " + total);
```


Clarifying elements through the pipeline

```
int total = IntStream.rangeClosed(1, 10)
    .filter(
        x -> {
            System.out.printf("%nFilter: %d%n", x);
            return x % 2 == 0;
        })
    .map(
        x -> {
            System.out.printf("map: %d", x);
            return x * 3;
        })
    .sum();
System.out.println("\n\nTotal: " + total);
```

```
Filter: 1
Filter: 2
map: 2
Filter: 3
Filter: 4
map: 4
Filter: 5
Filter: 6
map: 6
Filter: 7
Filter: 8
map: 8
Filter: 9
Filter: 10
map: 10
Total: 90
```

Collectors

- The terminal operation **collect()** combines the elements of a stream into a single object, such as a collection.
- There are many pre-defined collectors:
 - `Collectors.counting()`
 - `Collectors.joining()`
 - `Collectors.toList()`
 - `Collectors.groupingBy()`

Collectors

- The terminal operation **collect()** combines the elements of a stream into a single object, such as a collection.
 - There are many pre-defined collectors:
 - **Collectors.counting()**
 - `Collectors.joining()`
 - `Collectors.toList()`
 - `Collectors.groupingBy()`
- Returns the number of elements in the stream.

Collectors

- The terminal operation **collect()** combines the elements of a stream into a single object, such as a collection.
- There are many pre-defined collectors:
 - `Collectors.counting()`
 - **`Collectors.joining()`**

Joins the elements of the stream together into a String, with a specified delimiter
 - `Collectors.toList()`
 - `Collectors.groupingBy()`

Collectors

- The terminal operation **collect()** combines the elements of a stream into a single object, such as a collection.
- There are many pre-defined collectors:
 - `Collectors.counting()`
 - `Collectors.joining()`
 - **`Collectors.toList()`**
Puts the elements of the stream into a `List<>` and returns it.
 - `Collectors.groupingBy()`

Collectors

- The terminal operation **collect()** combines the elements of a stream into a single object, such as a collection.
- There are many pre-defined collectors:
 - `Collectors.counting()`
 - `Collectors.joining()`
 - `Collectors.toList()`
 - **`Collectors.groupingBy()`**

Groups the elements in the stream according to some parameter and returns a `HashMap` keyed by the “groupingBy” parameter.

Another terminal: `forEach()`

- `forEach()` applies the given method to each element of the stream.
- The method must receive one argument and return `void`.

reduce()

- Rather than using predefined reductions (`.sum()`, `.max()`, etc), we can write our own reduction.

```
int total = IntStream.rangeClosed(1, 10)
                      .reduce(1, (x, y) -> x * y);
```


reduce()

- Rather than using predefined reductions (`.sum()`, `.max()`, etc), we can write our own reduction.

```
int total = IntStream.rangeClosed(1, 10)
                      .reduce(1, (x, y) -> x * y);
```

The starting value.

This is the value for `reduce(0)`

reduce()

- Rather than using predefined reductions (`.sum()`, `.max()`, etc), we can write our own reduction.

```
int total = IntStream.rangeClosed(1, 10)
                      .reduce(1, (x, y) -> x * y);
```

The operation to perform.

Must take 2 parameters.

(Because it takes 2 params, we need to use the parens in the lambda)

Miscellaneous tid-bits

Producing a Stream from an Array

```
int total = IntStream.of(someInts)
                    .sum();
```

Producing a Stream from a Collection

```
List<String> strings = new ArrayList<>();  
strings.stream();
```

Creating a String from an Array

```
String out = IntStream.of(someInts)
                      .mapToObj (String::valueOf)
                      .collect(Collectors.joining(" "));
```

Here, the `mapToObj ()` operator is new.

It uses the specified method to convert the input element to a new type.

Using lines in a file as a stream

```
Files.lines(Paths.get("src/main/resources/PDPAssignment.csv"))
```

Returns a **Stream<String>**

...flatMap()?

```
Pattern splitAtSpaces = Pattern.compile("\\s+");  
String someStrings[] = {"one row", "some more words", "any other words", "and  
once upon a time"};  
  
Object list = Stream.of(someStrings)  
                    .map(line -> splitAtSpaces.splitAsStream(line))  
                    .collect(Collectors.toList());
```

What is the type of `list` after this is run?

How many elements are in the list?

4 elements in the final `list`.

(one for each entry in `someStrings`)

...flatMap()?

```
Pattern splitAtSpaces = Pattern.compile("\\s+");  
String someStrings[] = {"one row", "some more words", "any other words", "and  
once upon a time"};  
  
Object list = Stream.of(someStrings)  
                    .map(line -> splitAtSpaces.splitAsStream(line))  
                    .collect(Collectors.toList());
```

```
list => [{"one", "row"}, {"some", "more", "words"}, {"any", "other", "words"}, ...]
```

What is the type of `list` after this is run?

How many elements are in the list?

4 elements in the final `list`.

(one for each entry in `someStrings`)

...flatMap()?

```
Pattern splitAtSpaces = Pattern.compile("\\s+");  
String someStrings[] = {"one row", "some more words", "any other words", "and  
once upon a time"};  
  
Object list = Stream.of(someStrings)  
                    .flatMap(line -> splitAtSpaces.splitAsStream(line))  
                    .collect(Collectors.toList());
```

When I really want 13 items in the final list (one for every word in the original input), I use `flatMap()`.

When the output of a `map()` is a collection, `flatMap()` flattens the result by adding all the items in the output to the stream individually, rather than as a collection.

...flatMap()?

```
Pattern splitAtSpaces = Pattern.compile("\\s+");  
String someStrings[] = {"one row", "some more words", "any other words", "and  
once upon a time"};  
  
Object list = Stream.of(someStrings)  
                    .flatMap(line -> splitAtSpaces.splitAsStream(line))  
                    .collect(Collectors.toList());
```

```
list => ["one", "row", "some", "more", "words", "any", "other", "words", ...]
```

Immutability

- A tenet of functional programming is *immutability*
 - An object is not mutable– it can't change
 - Rather than change state (mutate it), create a new copy with the new state
 - Helps with concurrency

Applying of this to Objects,
not just primitive Types




Functional Programming in Java


- Stream that gets mapped, filtered, reduced, and collected... in some order.
 - Intermediate operations are not executed until a terminal operation is called.
- Lambdas: unnamed methods (functions) that can be applied to a stream
- Declarative vs. imperative

Backing up a bit...

What is a function?

- All elements in the domain are mapped to a value in the codomain
- There cannot exist elements in the domain w/ no image in the codomain
- A single element in the domain can only map to one element in the codomain





A function ***IS***.
It doesn't ***DO***.

Important Concepts in Functional Programming

- First class functions
- Anonymous functions
- Closures
- Currying
- Lazy evaluation
- Referential transparency

What does it mean to have no side effects?

- No mutation of variables
- No printing to the console or devices
- No writing to files, databases, networks, ...
- No exception throwing

(really, this is “no intentional side effects”)

Functional Programming Benefits

- Easier to reason about b/c they are deterministic
- Easier to test— no side effects
- More modular
 - Functions have an input and output; no side effects, concurrent modifications, etc...
- Composition and recombination easier
- Inherently thread-safe

Referential Transparency

Properties:

- Self-contained
- Deterministic
- Will never throw an Exception
- Won't cause conditions that cause other code to fail
- Won't hang due to an external device

Using substitution to reason about code

$3 * 2$	+	$4 * 5$
6	+	$4 * 5$
6	+	20

Using substitution to reason about code

```
public static void main(String[] args) {  
    int x = add(mult(2, 3), mult(4, 5));  
}
```

What happens if I instead call:

What happens if I instead call:

```
int x = (26);
```

```
public static int add(int a, int b) {  
    log(String.format("Returning %s as the result of %s + %s", a + b , a , b));  
    return a + b;  
}
```

```
public static int mult(int a, int b) {  
    return a * b;  
}
```

```
public static void log(String s) {  
    System.out.println(s);  
}
```

An example

```
public class DonutShop {  
    public static Donut buyDonut (CreditCard creditCard) {  
        Donut donut = new Donut();  
        creditCard.charge(Donut.price);  
        return donut;  
    }  
}
```

Can you see the side effect?

Create a Payment

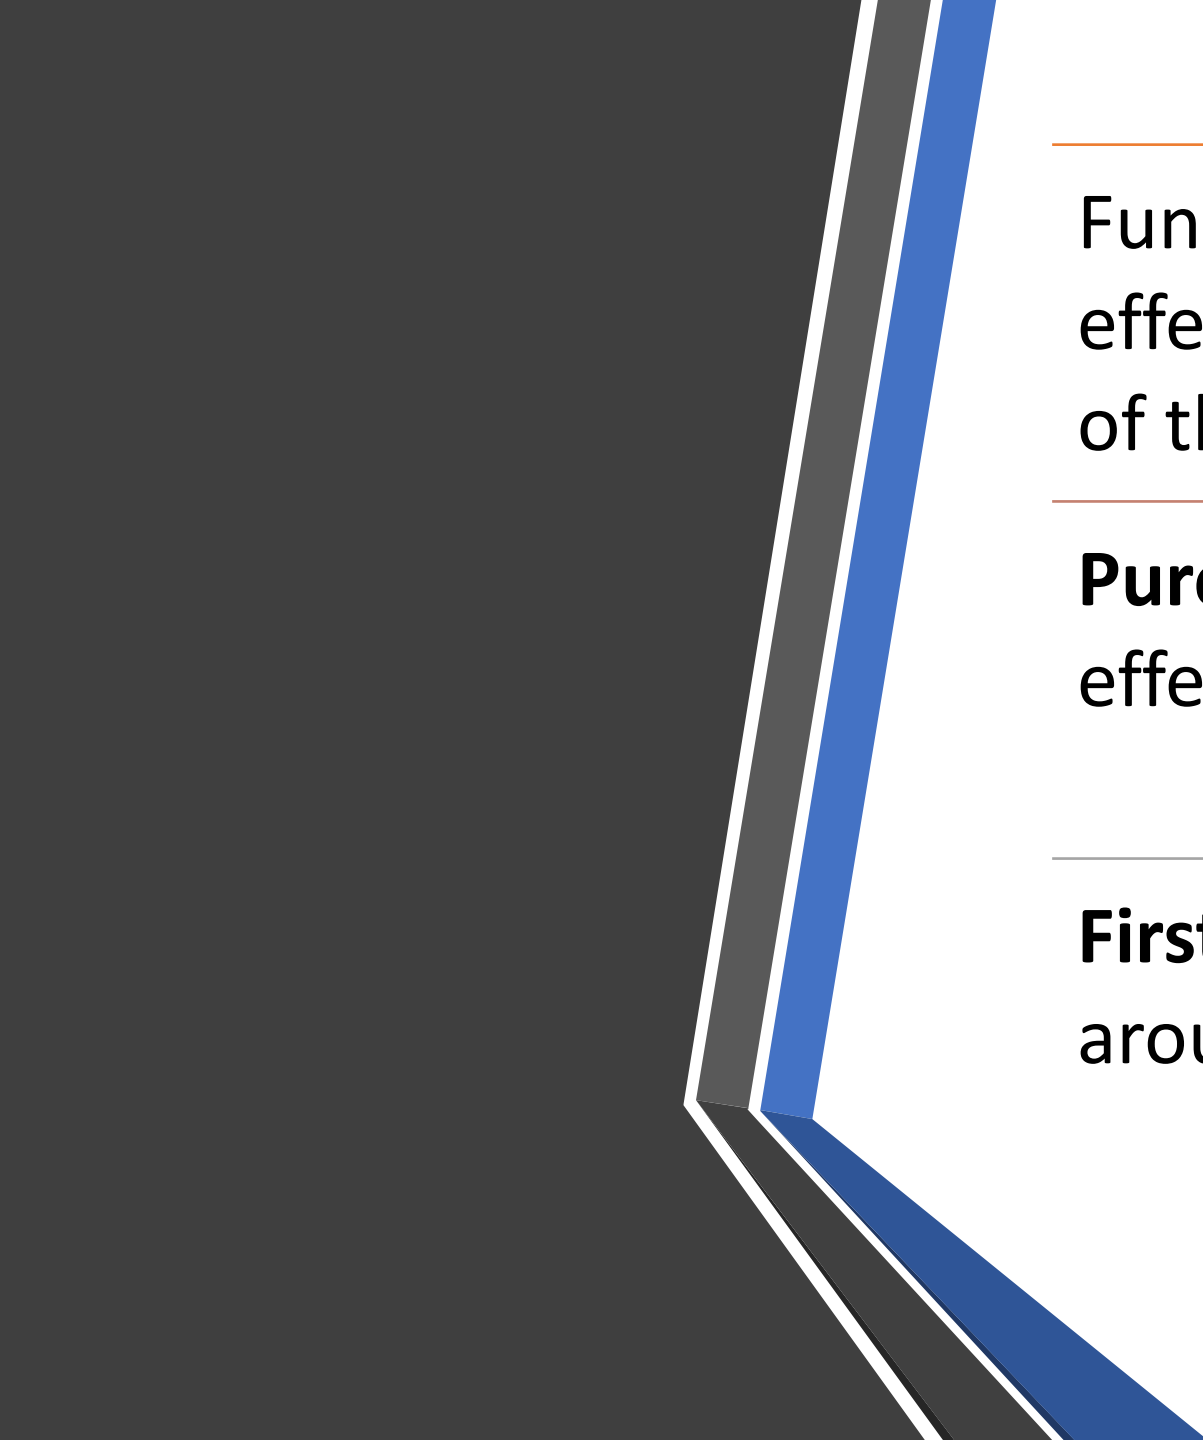
```
public class Payment {  
    public final CreditCard creditCard;  
    public final int amount;  
  
    public Payment(CreditCard card, int amt) {  
        this.creditCard = card;  
        this.amt = amount;  
    }  
}
```

Create a Purchase

```
public class Purchase {  
    public Donut donut;  
    public Payment payment;  
  
    public Purchase(Donut donut, Payment payment) {  
        this.donut = donut;  
        this.payment = payment;  
    }  
}
```

Modify buyDonut

```
public static Purchase buyDonut(CreditCard creditCard) {  
    Donut donut = new Donut();  
    Payment payment = new Payment(creditCard, Donut.price);  
    return new Purchase(donut, payment);  
}
```



Functional programming replaces side effects with returning a representation of the effects

Pure functions: functions without side effects

First-class functions: can be passed around as data/values.

Simplifying functions by Currying

- In some scenarios, you may want to simplify functions to only take a single parameter
- Some languages favor curried functions to achieve multiple parameters
 - ML and Haskell
 - Lambda calculus
- Simplifies your functions
 - Allows scalability

Currying

- Let's start with $f(x, y) = x + y$
 - This is really a function of 1 variable: $f((x, y))$
- What if I have $f(x)(y) = g(y)$?
 - $g(y) = x + y$
- Then, $f(x) = g$
 - The result of applying function f to x is a new function g
 - Then $g(y) = x + y$
- $f(x)(y)$ is the curried version of $f(x, y)$

Add 2 numbers

```
int add (int a, int b) {  
    return a + b;  
}
```

As a lambda:

$(a, b) \rightarrow a + b;$

Curried:

$(a) \rightarrow ((b) \rightarrow a + b);$

$a \rightarrow b \rightarrow a + b;$

Functional Interface: Function<T, R>

T – Type of the input to the function

R – Type of the result of the function

Modifier and Type	Method and Description
default <V> Function <T,V>	andThen (Function <? super R,? extends V> after) Returns a composed function that first applies this function to its input, and then applies the after function to the result.
R	apply (T t) Applies this function to the given argument.
default <V> Function <V,R>	compose (Function <? super V,? extends T> before) Returns a composed function that first applies the before function to its input, and then applies this function to the result.
static <T> Function <T,T>	identity () Returns a function that always returns its input argument.

Add 2 numbers

```
Function<Integer, Function<Integer, Integer>> add = (a) -> (b) -> a + b;
```

```
Int result = add
```

```
.apply(5)
```

```
.apply(6);
```

Functional in Java

A method is functional if:

- Does not mutate anything outside the function
- Does not mutate argument
- Does not throw errors or exceptions
- Always returns a value
- When called with the same argument, always returns the same result

Let's look at
some code





Questions?



Finishing up

- UML Design Draft due Monday
- Next Wednesday is our last class
 - Course wrap-up
 - Ask me (almost) anything (AMA)
- Homework #6 due **Sunday** @ 11:59pm

